

LEARNING UNIT 2 · ACTIVITY 2b

Client Rapport Assistant with Human Handover

The agent drafts. A licensed human approves. Nothing sends unsupervised.

PLATFORM

Microsoft Copilot Studio

APPROVAL

Microsoft Teams

DURATION

90 minutes

01



The Problem

Why Meridian cannot just automate this

SCENARIO

Meridian Asset Management today

Clients write in scared

Six discretionary portfolios. When markets move, clients email — worried about performance, volatility and NAV. They are anxious about their own money.

Replies take three days

The relationship managers are drowning. Some clients get a warm, careful answer. Some get a rushed one. Consistency is a matter of who was on duty.

One sentence is a breach

A manager under pressure once wrote “don’t worry, it always bounces back.” That is unlicensed financial advice in every jurisdiction that has a regulator.

The team wants AI to reply faster. Compliance says no. Both are right.



The AI is drafting a letter from a licensed firm to a retail investor about their money.

Getting the tone wrong is embarrassing. Getting the content wrong — “hold on, it will recover” — is an offence.

02



Human in the Loop

What the phrase actually means


DEFINITIONS

Three patterns, often confused

	HUMAN IN THE LOOP	ON / OUT OF THE LOOP
Who decides	AI proposes	AI decides
Who acts	Human authorises	AI acts
Can AI proceed alone?	No — it blocks	Yes
Supervision happens	Before the act	After, if at all
Failure mode	Delay	The client already read it

LU1's onboarding flow and LU2's investment advisor are both OUT of the loop. This lab is the only one genuinely IN it.

Where the loop closes



Everything before the gate is a proposal. Everything after it is a consequence. The node between them does nothing at all except wait.



The test of a real gate: submit an enquiry, then walk away.

The run says Running. It will still say Running tomorrow. Nothing times out, nothing defaults, nothing proceeds.

03

What You Build

Eleven nodes, one gate

The workflow on the canvas

[missing screenshot: lu2b-10-flow-canvas.png]

The eleven nodes, in order

1

When a HTTP request is received

Anyone (no auth) · relative path blank · six-field schema

2

Normalise_Enquiry (Compose)

Ticket ID, trimmed name, lowercased email — the safety control

3

Rapport_Agent (Agent)

No knowledge source · temperature 0.2 · structured output

4

Response

Five fields back to the browser. The draft is not one of them

5

Log_Draft (Excel)

Thirteen columns. The draft exists; no human has seen it.
Then four branch nodes: send + log on approve, assign + email on reject.

6

Human review

Assigned to a licensed person · delivered in Teams

04

Inside the Agent

What it may and may not write

The non-advisory rule

Never recommend

Buying, selling, holding, switching or redeeming anything

Never predict

Future returns, prices or NAV — no forecasts, no estimates

Never guarantee

“Markets always recover” · “you will not lose money”

Never judge suitability

Not safe, not risky, not right for them

Never time the market

No comment on whether now is a good or bad moment

Never invent a figure

No fee, NAV or return that was not in the client's own message

And what it MAY do: name the emotion, restate the concern, explain the process, offer a call with a licensed manager.



When in doubt, say less and offer the call.

That single sentence does more compliance work than the six prohibitions above it.

Six compliance flags

ADVICE _REQUE STED	asks what they should do	ESCALATE
GUARA NTEE_S OUGHT	wants a promise about returns	ESCALATE
VULNER ABLE_C LIENT	distress, illness, savings they cannot lose	ESCALATE
LEGAL_ OR_MED IA_THRE AT	mentions a lawyer, MAS, the press	ESCALATE
COMPL AINT	dissatisfied with Meridian or its fees	log only
WITHDR AWAL_I NTENT	raises redeeming or closing the account	log only

The four that escalate are the four that cannot be answered by a drafted email at all. For those, the correct output is a draft that says so.

The output contract

Custom structured output — not a Parse JSON node

```
{
  "ticketId":      "MAM-20260802-a607f8",
  "concernCategory": "Portfolio Performance",
  "emotionalTone":  "Concerned",
  "urgency":       "Medium",
  "complianceFlags": ["ADVICE_REQUESTED", "WITHDRAWAL_INTENT"],
  "escalate":      true,
  "suggestedSubject": "Your enquiry about recent performance (MAM-...)",
  "draftReply":     "<p>Thank you for writing, and I am sorry ...</p>"
}
```

emotionalTone is a FIELD, not a paragraph. That is what lets you sort a queue and show a manager that the angriest client has waited longest. Prose cannot be sorted.

Why structured output, not Parse JSON

With Text response you are asking and hoping

The model eventually wraps its answer in ```json fences, or opens with “Here is the JSON” — and the parse fails.

Structured output makes malformed JSON impossible

The schema is enforced at generation time. There is nothing to hope for.

Same discipline as the disclaimer

A rule the model physically cannot break beats a rule it is told not to break.

Reference it with a slash, not nested brackets

```
body('Rapport_Agent')?['structuredOutput/draftReply']
```

05



The Gate

Node 6 — Human review

What the approver is shown

[missing screenshot: lu2b-11-human-review.png]

The two inputs

Outcome · choice

Approve / Reject

Drives the If/Else branch.

⚠ Leave the DEFAULT BLANK.

Pre-filling it with 'Approve' turns the gate into a rubber stamp with a one-word setting nobody can see on the canvas.

If a default is required, use Reject —
then inattention fails safe.

Name · text

The approver types their own name.

Lands in the Approved By column of the audit table.

⚠ It is SELF-DECLARED. The node publishes no responder identity, so nothing stops someone typing 'Bob'.

A convention — where a captured identity would be evidence.

Approving in Microsoft Teams

[missing screenshot: lu2b-12-teams-approval.png]


06



The Client's View

What the browser is allowed to see

The Meridian client portal



A floating chat widget on a single-page asset-management site. The lab configuration panel holds the workflow URL.

Raising a concern

M

Meridian Asset Management
SINGAPORE

PortfoliosHow we respondLab setupTalk to us

Training environment. Meridian Asset Management is a fictitious fund manager used for this lab. Nothing here is financial advice, and no investment service is offered.

CLIENT RELATIONSHIP TEAM

Markets move. Your relationship manager does not disappear.

When your portfolio moves sharply, you want a person — not a form, and not a chatbot pretending to be one. Raise a concern here and a licensed relationship manager reads it, approves the reply personally, and comes back to you by email.

Raise a concern

Acknowledged instantly · Answered by a licensed representative

Our portfolios

Six discretionary mandates, from capital preservation to global growth. Past performance is not indicative of future performance.

“

My portfolio is down 14% this quarter and I do not understand why. Should I move?

AnxiousADVICE_REQUESTED

The assistant reads the concern and a human must handle, and drafts a response.

M

Meridian Client Care
A licensed manager reviews every reply

Good day. This is the Meridian client relationship team.

Tell us what is concerning you about your portfolio and we will get it in front of your relationship manager. We do not give investment advice through this chat.

Rachel Ongtrainer@meridian.exam

MAM-77401Global Growth

My portfolio is down 14% this quarter and I do not understand why. Should I move?

Meridian does not provide financial advice by chat. Messages are logged and reviewed by a licensed representative.

Name, email, account reference and portfolio — then the message.

The trainer dropdown fills all of it except the email, which it clears on purpose.

The approved reply is a REAL email. Always use your own address.

The receipt — and what is missing from it

The screenshot shows the Meridian Asset Management website. The header includes the company logo, navigation links (Portfolios, How we respond, Lab setup, Talk to us), and a disclaimer: "Training environment. Meridian Asset Management is a fictitious fund manager used for this lab. Nothing here is financial advice, and no investment service is offered." The main content area features a section titled "Markets move. Your relationship manager does not disappear." with a sub-header "CLIENT RELATIONSHIP TEAM". The text explains that when a portfolio moves sharply, a person (not a form or chatbot) should be contacted. A button "Raise a concern" is present, followed by the text "Acknowledged instantly · Answered by a licensed representative". Below this is a section "Our portfolios" with a brief description. A chatbot window is open, showing a user's message: "My portfolio is down 14% th understand why. Should I m". The chatbot responds with a message received notification, a reference number (MAM-20260802-a607f8), and a priority of Medium. The chatbot also states that a licensed relationship manager will review the reply personally and reply by email, and that investment advice will not be sent through the chat.

Reference, priority, and — because TC2 escalates — a line saying a manager will call.

NOT returned: the draft itself, emotionalTone, concernCategory, complianceFlags.

“We have classified you as Angry and flagged you as a COMPLAINT” is a disclosure no compliance officer would sign.



**Five fields go back to the browser.
The draft is not one of them.**

The agent's assessment of a person is an internal record. The browser is not where it belongs.

07



Run It

Eight test cases

The test matrix

TEST CASE	INPUT	EXPECTED
TC1 · Calm volatility	“dipped 6%, could someone explain?”	no flags → Approve
TC2 · Asks what to do	“Should I move everything to cash?”	ADVICE_REQUESTED
TC3 · Wants a guarantee	“promise I will not lose money”	GUARANTEE_SOUGHT
TC4 · Angry about fees	“furious · unacceptable · evasive”	COMPLAINT → Reject
TC5 · Wants to redeem	“redeem my entire holding”	WITHDRAWAL_INTENT
TC6 · Distressed retiree	“I am 68, cannot afford to lose this”	VULNERABLE → Reject
TC7 · Lawyer and MAS	“I have spoken to my lawyer”	LEGAL → Reject
TC8 · Factual NAV query	“how is the NAV calculated?”	no flags → Approve

Afterwards: Drafts has 8 rows, Approved_Replies has 5, Handover_Queue has 3.

Read the drafts, not just the flags

What a good TC2 draft does

Names the emotion in the first sentence.

Says plainly that the manager cannot give a recommendation by email.

Offers a call with a licensed person.

Contains no figure the client did not

supply.
This is the most valuable failure in the lab. Do not fix it quietly.

What a bad TC2 draft does

“Markets typically recover over the long term.”

You have just watched a language model commit a regulatory offence — in front of the whole room.

And the approval gate caught it.

That is what it is for.

Neither table has an empty accountability column

Approved_Replies names who authorised the email

That person is answerable for what reached the client.

Handover_Queue names who owes the client a call

A decline is a reassignment, not a deletion — the client is still waiting.

The regulator's question is not 'what did the machine decide'

It is 'which licensed person is answerable for it'.

Drafts records what the machine proposed, including rejected drafts

Asset or liability in litigation? Argue both sides in the debrief.

08

What Goes Wrong

Field-tested

The four that cost the most time

There is no workflow() function in this designer

Unknown function: workflow. Use guid() for the ticket ID instead.

The Instructions box escapes underscores

A pasted reference becomes Normalise_Enquiry, resolves to EMPTY, and the agent says 'no enquiry was included'. Always use the ⚡ picker.

The Agent node has no separate user-message field

The enquiry is appended to the END of the Instructions.

assignedTo must be a tenant user, picked from the dropdown

External addresses and tabbed-away text both fail at runtime with BadRequest.

Two that change behaviour silently

escalated must be UNQUOTED in the Response body

Quoted, it returns the string “false” — and `Boolean(“false”)` is true. Every calm client is told a manager is calling.

Compliance Flags must be wrapped in join()

It is an array. Without `join()` Excel writes `System.Object[]` and the audit column is useless.

Web search must be OFF

On, the agent can pull live market commentary into a client letter. One toggle undoes the entire non-advisory rule.

If the approval never arrives, switch Channel to Teams

The Outlook route silently fails to deliver on some tenants while the run still shows Running.

09



Debrief

The questions worth arguing about

DISCUSSION

Ask these before anyone packs up

1

TC6 is the hard one

Should a distressed 68-year-old's enquiry have reached the AI at all?

2

The approval button is a rubber stamp

After forty of these, what makes careful reading the path of least resistance?

3

emotionalTone is a judgement about a person, in a spreadsheet

Under the PDPA, is that personal data? Who can see it? How long is it kept?

4

Three controls, ranked

Which is probabilistic, which procedural, and which structural?

5

Approved By is self-declared

Is that an audit trail? What would close the gap?
It is not the technology, and it is not the chance of the model being wrong.

6

Compare with LU1 and Activity 2

Same platform, same shape, no gate. What is the actual variable?



The variable is not how likely the model is to be wrong.

It is what a wrong sentence costs — and to whom.

Generalities to the public can go unsupervised. This client's money cannot.

10



Going Further

Where this goes next

If you had another day

Make rejection informative

Add a required Comments input on reject. Right now Handover_Queue records that a draft was refused but never why — every Status cell says the same sentence.

Route before the model sees it

A keyword pre-filter that sends distress straight to a human, with no draft at all. What does it cost, and what does it buy?

Close the identity gap

The node will not tell you who approved. What would you have to build to capture it — and is an audit trail you must remember to maintain an audit trail at all?